



HTOC

HEALTHCARE THREAT OPERATIONS CENTER

Next Observed Indicator Forecasting — NextObservedIndicatorV3.0

Version 3.0
April 2026

Table of Contents

Table of Contents

Table of Contents	2
1. Purpose	4
2. Scope	4
3. Prerequisites	4
3.1 Environment	4
3.2 Input Data Files	5
4. Key Configuration Parameters	5
5. Pipeline Overview	6
6. Step-by-Step Execution	7
Step 1 — Import Dependencies	7
Step 2 — Set Date Window	7
Step 3 — Load Daily OpDiv CSVs	8
Step 4 — Clean and Reshape	8
Step 5 — Build Per-OpDiv Dense Panel	9
Step 6 — Define Modeling Functions	9
Step 7 — Run <code>main()</code>	10
7. Features Used in Modeling	11
7.1 Time-Series Features (computed per indicator)	11
7.2 Supervised Labels (computed per indicator per horizon)	12
8. Forecasting Models	12
8.1 Logistic Regression	12
8.2 Gradient Boosting Classifier	13
8.3 Exponential (Poisson-Rate) Model	13
8.4 Weibull AFT (Accelerated Failure Time) Model	13
8.5 Weighted Ensemble	14
9. Confidence Labels	14
10. Production Output Table	14
11. How to Run the Notebook	15
11.1 Open the Notebook	15
11.2 Verify Prerequisites	16
11.3 Run All Cells	16

11.4 Monitor Progress	16
11.5 Enable File Output (Optional)	17
11.6 Typical Run Time	17
11.7 Scheduled Execution	18
12. Troubleshooting	18
13. Optional: Feedback Loop and Model Retraining	19
How it works:	19
Notes:	20
14. Interpreting Results	20
15. Maintenance Notes	21
16. Appendix - Standalone Python Script	21
17. Related Documents	22

1. Purpose

This SOP describes how to operate the Next Observed Indicator (NOI) forecasting pipeline. The notebook loads daily OpDiv observation records, engineers time-series features for each indicator per OpDiv, and runs a multi-model ensemble to forecast the probability that each indicator will be observed again within **1, 7, 14, 30, and 45-day** horizons. Results are presented as a production-ready prioritization table with categorical confidence labels.

-

2. Scope

This procedure applies to HTOC analysts and data engineers who run, maintain, or consume indicator observation forecasts. It covers end-to-end execution of NextObservedIndicatorV3.0.ipynb, from environment setup through production output review. The optional feedback/retraining section is also documented for users who wish to improve model accuracy over time.

-

3. Prerequisites

3.1 Environment

Requirement	Notes
Python 3.x	Compatible with numpy, pandas, scikit-learn, lifelines

Requirement	Notes
Jupyter (VS Code or JupyterLab)	Notebook must be run interactively
scikit-learn	`LogisticRegression`, `GradientBoostingClassifier`, `StandardScaler`, `Pipeline`
lifelines	`WeibullAFTFitter` for survival-model probabilities
numpy, pandas	Standard data processing
Access to `Z:\HTOC\Data_Analytics\`	All input OpDiv CSV files live here

Install missing packages if needed:

```
pip install scikit-learn lifelines
```

3.2 Input Data Files

The notebook loads daily OpDiv observation CSVs using the following path template:

```
Z:\HTOC\Data_Analytics\Data\OpDiv_Observations\htoc_opdiv_obs_d{YYYYMMDD}.csv
```

Files must be available for the **100 days prior to today**. If any daily files are missing, the loader skips them gracefully, but gaps will reduce feature quality for affected indicators.

-

4. Key Configuration Parameters

These values are defined near the top of the notebook. Review them before each run.

Parameter	Default	Description
`start_date` offset	100 days before today	How far back daily CSV files are loaded
`end_date` offset	0 days (today)	End of the observation window
`n_days` (in `load_data`)	100	Used by the standalone `load_data` helper (not called in the main pipeline)
`high_thresh`	0.60	Minimum ensemble probability for a "Highly likely" confidence label
Secondary probability threshold	0.07	Minimum probability for "Possibly active" label
Frequency threshold (High)	≥ 2	Minimum observation frequency count for "Highly likely" classification
Frequency threshold (Possibly active)	≥ 1	Minimum observation frequency count for "Possibly active" classification
Exponential rate floor	1e-6	Prevents division by zero in the Poisson-style model

-

5. Pipeline Overview

The notebook executes in the following logical stages. Cells must be run in order from top to bottom.

```
[1] Import Dependencies
    ↓
[2] Set Date Window & File Path Template
    ↓
[3] Load & Concatenate Daily OpDiv CSVs
    ↓
[4] Clean & Reshape Data
    ↓
[5] Build Per-OpDiv Dense Panel (Cartesian Product)
    ↓
[6] Define Modeling Functions
    ↓
[7] Run main() – Feature Engineering + All Models + Ensemble
    ↓
[8] Display Production Output Table
    ↓
[9] (Optional) Feedback Loop & Model Retraining
```

-

6. Step-by-Step Execution

Step 1 — Import Dependencies

The first cell imports:

```
numpy, pandas, os, datetime, warnings
```

```
sklearn: LogisticRegression, GradientBoostingClassifier, StandardScaler, Pipeline
```

```
lifelines: WeibullAFTFitter
```

Verify: No import errors. If `lifelines` is missing, install it before proceeding.

-

Step 2 — Set Date Window

The notebook calculates:

```
start_date = today - 100 days
```

```
end_date = today
```

A list of all dates in this range is generated (`datelist`) and printed. The file path template for OpDiv CSVs is also defined here.

Verify: The printed date list covers the expected 100-day window ending today.

-

Step 3 — Load Daily OpDiv CSVs

The `load_files()` function iterates over `datelist`, attempts to read each daily CSV at the defined path template, skips missing files silently, and concatenates all loaded files into a single dataframe `src`.

Verify: `src` is non-empty and contains the expected columns (`indicator`, `OpDiv`, `obs_date`, `observations`). If `src` is empty, check that the file path template and [Z:\](#) drive mapping are correct.

-

Step 4 — Clean and Reshape

Drops `curr_date` and `indicator_key` columns if present.

Renames `obs_date` → `date`.

Parses `date` as `datetime`.

Strips whitespace from `indicator` and `OpDiv` columns.

Step 5 — Build Per-OpDiv Dense Panel

For each OpDiv in the data, a **full Cartesian product** is created across:

All unique API users

All calendar days from the minimum date in data through today

All unique indicators

This dense grid is left-merged with actual observation counts. Missing days are filled with 0 observations. Additional calendar features are added:

Column	Description
`dayofweek`	Integer 0–6 (Monday = 0)
`is_weekend`	1 if Saturday or Sunday, else 0
`day`	Day of month
`month`	Month number
`seen`	Binary: 1 if `observations > 0`, else 0

The result is stored in `opdiv_merged`, a dictionary keyed by OpDiv name.

Step 6 — Define Modeling Functions

This cell defines all helper functions used by `main()`. No computation is performed. Key functions:

Function	Purpose
<code>`extract_time_series_features(group)`</code>	Computes per-indicator features from the <code>`seen`</code> series
<code>`build_features(df)`</code>	Applies <code>`extract_time_series_features`</code> across all indicators for an OpDiv
<code>`train_predict(X, y)`</code>	Trains and scores a Logistic Regression pipeline
<code>`train_gbt(X, y)`</code>	Trains and scores a Gradient Boosting Classifier
<code>`get_model_outputs(df)`</code>	Runs all four model types and merges results
<code>`add_rule_and_ensemble(df)`</code>	Applies rule-based labels and computes weighted ensemble
<code>`classify_window(prob, freq)`</code>	Returns confidence label string
<code>`add_confidence_and_format(df)`</code>	Formats probabilities as percentages and assigns confidence labels
<code>`build_production_output(df)`</code>	Selects and renames columns for the final output table

-

Step 7 — Run `main()`

The `main()` function is the core of the pipeline. It loops over each OpDiv in `opdiv_merged` and runs the full modeling sequence:

Calls `build_features()` to extract time-series features per indicator.

Calls `get_model_outputs()` to generate raw model probabilities.

Calls `add_rule_and_ensemble()` to apply rules and compute ensemble scores.

Calls `add_confidence_and_format()` to apply confidence thresholds and format output.

Calls `build_production_output()` to produce the final display table.

Results are stored in a dictionary (`outputs`) keyed by OpDiv name.

Prediction output path (currently commented out in the notebook):

`C:\Users\jaskew\Documents\NOI_Logs\Predictions\`

Uncomment the CSV save block inside `main()` if persistent file output is needed.

-

7. Features Used in Modeling

7.1 Time-Series Features (computed per indicator)

Feature	Description
<code>`last_seen`</code>	Days since the indicator was last observed
<code>`freq_1`</code>	Observation count in the last 1 day
<code>`freq_7`</code>	Observation count in the last 7 days
<code>`freq_14`</code>	Observation count in the last 14 days
<code>`freq_30`</code>	Observation count in the last 30 days
<code>`freq_45`</code>	Observation count in the last 45 days
<code>`avg_gap`</code>	Average number of days between observations

Feature	Description
`burstiness`	Coefficient of variation of inter-observation gaps (high = irregular)

7.2 Supervised Labels (computed per indicator per horizon)

Label	Description
`label_7`	1 if indicator was seen at least once in the last 7 days
`label_14`	1 if indicator was seen at least once in the last 14 days
`label_30`	1 if indicator was seen at least once in the last 30 days
`label_45`	Binary flag derived from the 45th-from-last row (see note below)

> **Note:** `label_45` uses `series[-45]` (a single index) rather than a slice `series[-45:]`. This means the 45-day label reflects only whether the indicator was seen on that specific day, not whether it was seen at all in the 45-day window. Be aware of this when interpreting 45-day probabilities.

-

8. Forecasting Models

The pipeline uses four models, combined into a weighted ensemble.

8.1 Logistic Regression

Implementation: `sklearn.pipeline.Pipeline` with `StandardScaler` + `LogisticRegression`

Horizons: 7, 14, 30, 45 days (1-day reuses the 7-day label as its training target)

Features: All 8 time-series features

Guard: If only one class is present in training data, the model returns NaN probabilities (handled downstream)

8.2 Gradient Boosting Classifier

Implementation: `sklearn.ensemble.GradientBoostingClassifier` (default hyperparameters)

Horizons: 7, 14, 30, 45 days (same 1-day reuse as logistic)

Features: All 8 time-series features

Purpose: Captures non-linear patterns and feature interactions

8.3 Exponential (Poisson-Rate) Model

Formula: $P(\text{seen within } t \text{ days}) = 1 - \exp(-\text{rate} \times t)$

Rate: `freq_30 / 30`, clipped to a minimum of $1e-6$

Horizons: 1, 7, 14, 30, 45 days

Purpose: Models indicators as a memoryless Poisson process; good for regularly-seen indicators

8.4 Weibull AFT (Accelerated Failure Time) Model

Implementation: `lifelines.WeibullAFTFitter`

Duration variable: `avg_gap`

Event variable: `label_7` (7-day seen label)

Horizons: Survival function evaluated at 1, 7, 14, 30, 45 days

Purpose: Models time-to-next-observation, capturing "burstiness" and irregular patterns

8.5 Weighted Ensemble

Final ensemble probabilities are computed as:

$$\text{ensemble} = 0.30 \times \text{rule_logistic} + 0.25 \times \text{GBT} + 0.25 \times \text{Weibull} + 0.20 \times \text{exponential}$$

Applied for each forecast horizon (1d, 7d, 14d, 30d, 45d).

-

9. Confidence Labels

Ensemble probabilities are classified into three confidence tiers:

Label	Condition
Highly likely	<code>`ensemble_prob >= 0.60` AND `freq >= 2`</code>
Possibly active	<code>`ensemble_prob >= 0.07` AND `freq >= 1`</code>
Low confidence	All other cases

Probabilities are formatted as percentage strings (e.g., "73%") in the output table.

-

10. Production Output Table

The final output table contains one row per indicator per OpDiv and includes:

Column	Description
Indicator	The threat indicator value
OpDiv	The observing organization division
Last Seen (days ago)	Days since most recent observation
Probability: 1-Day	Ensemble forecast for next-day observation
Probability: 7-Day	Ensemble forecast for 7-day window
Probability: 14-Day	Ensemble forecast for 14-day window
Probability: 30-Day	Ensemble forecast for 30-day window
Probability: 45-Day	Ensemble forecast for 45-day window
Confidence	Categorical label: Highly likely / Possibly active / Low confidence
Seen Today	Whether the indicator was observed on the most recent date in data

> **Known display note:** The `ensemble_45d` column may appear with its raw column name rather than the formatted "Probability: 45-Day" label in some runs, due to a rename mapping gap in `build_production_output`. This does not affect the underlying data.

-

11. How to Run the Notebook

11.1 Open the Notebook

Open **Cursor** (or VS Code / JupyterLab).

Navigate to <H:\HTOC\notebooks\observationEventForecasting\>.

Open NextObservedIndicatorV3.0.ipynb.

11.2 Verify Prerequisites

Before executing any cells, confirm the following:

- [] [Z:\HTOC\Data Analytics](#) is accessible and mapped.
- [] Daily OpDiv CSV files are present for at least the past 100 days in [Z:\HTOC\Data Analytics\Data\OpDiv Observations](#).
- [] `scikit-learn` and `lifelines` are installed in the active Python environment.

11.3 Run All Cells

Use **Run All** to execute the notebook from top to bottom:

Cursor / VS Code: Click the ►► Run All button in the notebook toolbar, or open the Command Palette (Ctrl+Shift+P) and select **Notebook: Run All Cells**.

JupyterLab: Kernel → Restart Kernel and Run All Cells... → confirm restart.

> **Important:** Cells must be run in sequential order. Each stage depends on dataframe state built by the previous stage.

11.4 Monitor Progress

Stage	Expected Output
Date window	Printed date list and today's date
CSV load	`src` dataframe displayed with raw observation records

Stage	Expected Output
Panel construction	<code>`opdiv_merged['DHA']`</code> (or another OpDiv) displayed as a sample dense panel
<code>`main()`</code>	Per-OpDiv production tables printed; <code>`display(production_output.head(5))`</code>

If a cell raises an exception, stop and do not run subsequent cells. Refer to the **Troubleshooting** section (Section 12) before restarting.

11.5 Enable File Output (Optional)

CSV output is currently commented out inside `main()`. To save per-OpDiv prediction files, uncomment these lines:

```
opdiv_output_dir = r'C:\Users\jaskew\Documents\NOI Logs\Predictions'
os.makedirs(opdiv_output_dir, exist_ok=True)
# ... save CSV block
```

Ensure the target directory exists or can be created.

11.6 Typical Run Time

Phase	Approximate Duration
CSV load (100 days of files)	1–3 minutes
Panel construction (Cartesian)	2–5 minutes (depends on number of OpDivs and indicators)
Feature engineering + modeling	3–10 minutes
Output formatting	< 30 seconds
Total	**~6–18 minutes**

11.7 Scheduled Execution

This pipeline currently runs **automatically every day at 7:00 AM** via **Windows Task Scheduler** on **F.R.E.D.** Manual execution (Sections 11.1–11.5) is only necessary when:

Troubleshooting a failed or incomplete scheduled run

Running an ad hoc forecast outside the normal schedule

Testing notebook changes before deploying them to the scheduled task

If the scheduled run fails, check the Task Scheduler history on F.R.E.D for the last execution status and any error output logged by the task. Confirm that F.R.E.D has access to the [Z:\](#) drive and that the daily OpDiv CSV files for the expected date were available at 7:00 AM when the task executed.

-

12. Troubleshooting

Symptom	Likely Cause	Resolution
`src` is empty after load	No CSV files found at the path template	Verify `Z:\` drive is mapped; check date range and file naming convention
`ImportError: lifelines`	Package not installed	Run `pip install lifelines` in the kernel
`NaN` probabilities in output	Single-class indicator (only ever seen or never seen)	Expected behavior; the model guard returns NaN and these are handled downstream

Symptom	Likely Cause	Resolution
Weibull AFT convergence warning	Sparse data for some indicators	Non-blocking; model still produces a result; consider excluding very-rare indicators
`ensemble_45d` appears as raw column name	Known `rename` gap in `build_production_output`	Cosmetic only; data is correct; rename column manually if needed for export
Very slow panel construction	Large number of OpDivs, users, and indicators	Reduce `start_date` offset (e.g., 60 days instead of 100) to shrink the panel
`label_45` interpretation unexpected	Single-index vs. window-slice bug in `extract_time_series_features`	Treat 45-day forecasts as approximate; prioritize 7/14/30-day horizons for decisions

-

13. Optional: Feedback Loop and Model Retraining

The last cell of the notebook contains a stub for a per-OpDiv feedback and retraining loop. This is **not run as part of the standard pipeline** but can be used to improve model accuracy over time.

How it works:

Load forecast logs from [Logs/<OpDiv>/forecast_log.xlsx](#) — these should contain previous forecasts with actual outcomes appended.

Merge logs into a `train_master_{OpDiv}.csv` file (deduplicated by indicator + date key).

Retrain a `GradientBoostingClassifier(n_estimators=200, max_depth=3)` on the merged training master.

Save the trained model to `gbc_7d_{OpDiv}.pk1` for future use.

Notes:

Forecast logs must be manually maintained or generated by a separate logging process; the main pipeline does not write them automatically.

A commented-out `classification_report` block is included for optional hold-out validation.

Integration of saved `.pk1` models back into the main `main()` pipeline requires additional development.

-

14. Interpreting Results

Confidence	Recommended Action
Highly likely	Prioritize for active monitoring; high probability of re-observation within the forecast window
Possibly active	Include in standard monitoring queue; moderate probability warrants continued tracking
Low confidence	Indicator is rare or irregular; monitor passively; re-evaluate if observed again

Use the **1-day probability** for immediate triage decisions. Use **7-day and 30-day probabilities** for workload planning and longer-horizon threat tracking.

High burstiness combined with a high probability score indicates an indicator that appears in clusters — pay attention to the "Seen Today" flag for these, as a fresh observation may signal an active campaign.

-

15. Maintenance Notes

No persistent model files are produced by the standard run. Models are trained in-memory each execution on the current 100-day window. This ensures forecasts always reflect recent behavior.

Extending horizons: Adding a new forecast horizon (e.g., 60 days) requires adding a `label_60` in `extract_time_series_features`, new model calls in `get_model_outputs`, updated ensemble weights, and a new `classify_window` call in `add_confidence_and_format`.

Ensemble weight tuning: Weights (`0.30 / 0.25 / 0.25 / 0.20`) are defined in `add_rule_and_ensemble`. Adjustments should be validated against historical forecast accuracy before deploying to production.

1-day forecast note: The 1-day logistic and GBT models are currently trained on the 7-day label (`y_7`), not a dedicated 1-day label. This is a known limitation — interpret 1-day ensemble probabilities with the understanding that the supervised component reflects 7-day patterns, partially corrected by the exponential and Weibull components.

-

16. Appendix - Standalone Python Script

The complete standalone script extracted from the notebook is attached at:

H:\HTOC\documentation\SOP\Appendix_Scripts\NextObservedIndicatorV3.0_standalone.py

Run with:

```
&"C:\Program Files\Python313\python.exe" "H:\HTOC\documentation\SOP\Appendix  
Scripts\NextObservedIndicatorV3.0_standalone.py"
```

-

17. Related Documents

HTOC OpDiv Observation Feed runbook

SOP_PRISM_ThreatAssessmentScoring.md — companion SOP for indicator risk scoring

[Z:\HTOC\Data Analytics\Data\OpDiv Observations\](#) (input data location)

Previous notebook versions: NextObservedIndicatorV1.0.ipynb, NextObservedIndicatorV2.0.ipynb